

1 Usage and structure of continuous integration in
2 open source projects

3 Joseph Ling
j1653@kent.ac.uk



School of Computing
University of Kent
United Kingdom

Word Count: 6,100

4 February 22, 2020

6 This paper describes a simple heuristic approach to solving large-scale con-
7 straint satisfaction and scheduling problems. In this approach one starts
8 with an inconsistent assignment for a set of variables and searches through
9 the space of possible repairs. The search can be guided by a value-ordering
10 heuristic, the *min-conflicts heuristic*, that attempts to minimize the num-
11 ber of constraint violations after each step. The heuristic can be used with
12 a variety of different search strategies. We demonstrate empirically that on
13 the n -queens problem, a technique based on this approach performs orders of
14 magnitude better than traditional backtracking techniques. We also describe
15 a scheduling application where the approach has been used successfully. A
16 theoretical analysis is presented both to explain why this method works well
17 on certain types of problems and to predict when it is likely to be most
18 effective.

1 Introduction

<https://arxiv.org/ftp/arxiv/papers/1703/1703.07019.pdf>

Continuous integration (CI) is becoming more popular over the last few years. This can be seen by how major version control hosting services Github, Bitbucket and Gitlab have all started to or have been improving their CI product. In terms of research, configuration as code Rahman, Mahdavi-Hezaveh and Williams (2019) and continuous integration Copeland (2010) with Shahin, Ali Babar and Zhu (2017) demonstrating breadth of the research.

Continuous integration is a process of automatically running compiling, running tests and checking that the product works. This can be combined with Continuous Delivery where the product is deployed or released after it has gone through CI.

This can get complicated quickly therefore configuration as code (or infrastructure as code) is used to configure it. The main kind of configuration format used for this is yaml (reference to what it is??) followed by xml and java based scripting formats.

In terms of looking at usage we are going to do a similar look at the data as did Michael Hilton, Marinov and Dig (2016). The important aspect will be looking at how usage has changed over the last 5 years along with looking more closely at which repositories are more likely to use CI/CD. For this we are going to focus on the following research questions: - usage of CI vs non usage - multiple CI used - per language CI usage - stars, subscribers and commits for likelihood of using CI

This should give us a better understanding of the sample of repositories from Github. From there we look at the structure of the configuration files to understand how certain aspects of it are used. - configuration errors when loading the config (just yaml parsing errors atm) - how are comments used in the configuration? - how scripts with the configuration files? (need to elaborate more on this one)

A key aspect is that these questions do not look too deeply into the

individual implementation of each CI system. This is because there are already some good papers looking Gallaba and McIntosh (2018) at this but in order to be able to compare the different configuration types it is important to compare similar attributes (there is also a time factor in here as well).

need to carefully plan out how to gage the size of repo, as I don't want to have too download repos

2 Previous Works

Configuration as code or infrastructure as code has been an increasing area of research over the last few years. There seems to be slightly more research in infrastructure as code Rahman, Mahdavi-Hezaveh and Williams (2019) . The has been a focus on Puppet and Chef, for example in Sharma, Fragkoulis and Spinellis (2016) looks at code quality by the measure of "code smell" of Puppet code. This tackles the problem by defining by best practices and analyzing the code against that. In the case of Cito et al. (2017) it uses the docker linter in order to be able to analyse the files.

Config as code is not necessarily infrastructure as code is slightly differently. So the comparison isn't fair necessarily or accurate.

15, 18, 23, 30

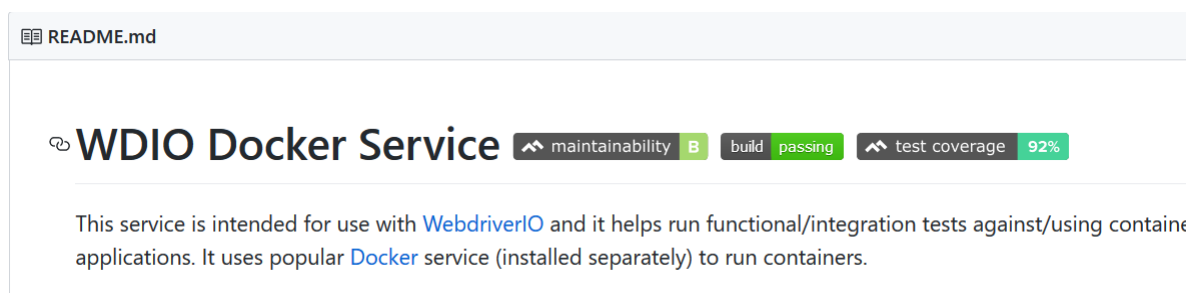
docker, puppet, travis analysis

3 Methodology

In order to get repositories with CI/CD configuration from Github we have a number of approaches. The first is too use the search for particular files but this is limited to only 1000 results. The alternative is to search for repositories and we bypass the 1000 result limit to an extent by getting results for every 'star' count (stars are used to like or upvote a repository). Although this will be giving us a lot of results it will still only be a sample of the population but will give us a wider range of results. As their is rate limiting multiple github api keys can be used to speed up the scraping of data (gTorrent could also be used to speed up the process I think).

After we have got a repository we need to get the CI/CD files from it. This is fairly easy as the CI/CD systems normally require a strict naming

79 convention and location within the repository. However as most of them are
80 yaml based you can have ".yaml" and ".yml" and users can use all sorts of
81 mixtures of upper and lower case. We try to account for this but won't get
82 every scenario. This combined with the fact that we are only looking for
83 top configuration files based on github (2017) along with github actions and
84 azure pipelines. Is why we also check repositories for their ReadMe.md file
85 to check if it has a build tag. sdfghjkl



86

where did this image come from??
reference it man

87 fghjk In doing so it should give a wider net when sampling and help to un-
88 derstand when a CI system is either not using configuration as code or using
89 a different CI system.

90 Results and the spread of the actual data with the watches vs stars maybe
91 the star count over time to demonstrate the floors in the data gathering
92 process

93 There are dangers in scraping data off github in terms of assumptions to
94 do with the population as found in Kalliamvakou et al. (2014). In Github
95 you can fork a repository which copies in order to remove these we check for
96 fork flag on the repository. This causes are dataset to go from NUMBER to
97 OTHER_NUMBER.

98 Additionally the assumption that all repositories are of programming
99 projects with code in them is wrong. A number of repositories can be used
100 for storage, experimental, academic and other things. However they to all
101 some extent can use CI/CD for their work as a number of books were found
102 when looking through the dataset could use CI/CD.

Analyse of readme
can be used to try
and classify but I
don't think that is
necessary. I think
the key factor is
that any concluding
remarks anywhere
take this into ac-
count.

104 3.1 CI/CD spread in the sample

105 This leads us to get the following data:


| CI/CD | count | configs per repo | duplicates | duplicate percent |
|-----------------|-------|------------------|------------|-------------------|
| config file | 8327 | 39% | 1221 | 15% |
| found in ReadMe | 582 | 3% | | |
| none found | 11469 | 53% | | |

106 A repository on github is like a folder so can contain any number of
107 configuration files in it. Therefore we can get any number of configuration
108 files in that folder. This is taken into account with the second pair of columns
109 for the first row. It demonstrates that there are a large number of repositories
110 with multiple kinds of configuration (todo make sure that github actions
111 multiple file thing isn't calculated here).


112 The next row is for when we couldn't pick up the configuration used for
113 CI/CD and check the ReadMe.md file for build status tag.

114 The final row is shows the repositories that either don't have any config-
115 uration or no configuration that could be found.

116 However that doesn't give us too much insight into the dataset. Here is a
117 graph showing the subscribers plotted against the number of stars. The key
118 here to understand is not potentially any correlation but to see the spread
119 of data that the table is showing.



121 The second graph helps give a understanding to the give a depth of the
122 data for where the graph is just blue. This is because on Github you get
123 more repositories with smaller star counts than large ones.



125 - what is the spread of ci/cd systems for public github repositories? this
126 will take into account operating system, programming language, star count,
127 subscriber count - note: something along the lines of multiple configuration
128 files - what naming convention do they use for the files? (in order to under-
129 stand common practices)

130 this will follow on from the previous graphs looking at spread of CI/CD
131 configs found in the whole sample

132 then look at the difference that large repositories more than 100 commits
133 with more than 2 contributors

134 then look at recent commits

135 perhaps a small look at naming conventions used???

136 4 Config file results

137 4.1 configuration errors when loading the config (just 138 yaml parsing errors atm)

| | config | percentage |
|-----------------|--------|------------|
| travis | 7051 | 74% |
| github | 1544 | 16% |
| circleci | 759 | 8% |
| jenkinsPipeline | 113 | 1% |
| drone | 54 | 1% |
| buildkite | 20 | 0% |
| teamcity | 4 | 0% |
| azure | 1 | 0% |
| semaphore | 1 | 0% |

| yaml_encoding_error | composer error | constructor error | parse error | scanner error |
|---------------------|----------------|-------------------|-------------|---------------|
| circleci | 1 | 0 | 0 | 1 |
| drone | 20 | 0 | 0 | 0 |
| github | 0 | 1 | 0 | 2 |
| travis | 4 | 0 | 6 | 16 |

139 **4.2 How are comments used in configuration?**

140 **4.3 How are stages used in configuration?**

141 and looking into branches

142 **4.4 How are script tags used?**

143 - how scripts with the configuration files? (need to elaborate more on this
144 one)

145 By almost any measure, the Hubble Space Telescope scheduling problem
146 Between ten thousand and thirty thousand astronomical observations per
147 year must be scheduled, subject to a great variety of constraints including
148 power restrictions, observation priorities, time-dependent orbital character-
149 istics, movement of astronomical bodies, stray light sources, etc. Because the
150 telescope is an extremely valuable resource with a limited lifetime, efficient
151 scheduling is a critical concern. An initial scheduling system, developed us-
152 ing traditional programming methods, highlighted the difficulty of the prob-
153 lem; it was estimated that it would take over three weeks for the system to
154 schedule one week of observations. As described in section 6, this problem
155 was remedied by the development of a successful constraint-based system to
156 augment the initial system. At the heart of the constraint-based system is
157 a neural network developed by Adorf and Johnston, the Guarded Discrete
158 Stochastic (GDS) network,

159 From a computational point of view the network is interesting because
160 Adorf and Johnston found that it performs well on a variety of tasks, in

161 addition to the space telescope scheduling problem. For example, the network
 162 performs significantly better on the n -queens problem than methods that
 163 were previously developed. The n -queens problem requires placing n queens
 164 on an $n \times n$ chessboard so that no two queens share a row, column or diagonal.
 165 The network has been used to solve problems of up to 1024 queens, whereas
 166 most heuristic backtracking methods encounter difficulties with problems
 167 one-tenth

168 In a standard Hopfield network, all connections between neurons are sym-
 169 metric. In the GDS network, the main network is coupled asymmetrically to
 170 an auxiliary network of *guard neurons* which restricts the configurations that
 171 the network can assume. This modification enables the network to rapidly
 172 find a solution for many problems, even when the network is simulated on
 173 a serial machine. Unfortunately, convergence to a stable configuration is no
 174 longer guaranteed. Thus the network can fall into a local minimum involving
 175 a group of unstable states among which it will oscillate. In practice, however,
 176 if the network fails to converge after some number of neuron state transitions,
 177 it can simply be stopped and started over.

178 To illustrate the network architecture and updating scheme, let us con-
 179 sider how the network is used to solve binary constraint satisfaction problems.
 180 A problem consists of n variables, $X_1 \dots X_n$, with domains $D_1 \dots D_n$, and a
 181 set of binary constraints. Each constraint $C_\alpha(X_j, X_k)$ is a subset of $D_j \times D_k$
 182 specifying incompatible values for a pair of variables. The goal is to find
 183 an assignment for each of the variables which satisfies the constraints. (In
 184 this paper we only consider the task of finding a single solution, rather than
 185 that of finding all solutions.) To solve a CSP using the network, each vari-
 186 able is represented by a separate set of neurons, one neuron for each of the
 187 variable's possible values. Each neuron is either "on" or "off", and in a solu-
 188 tion state, every variable will have exactly one of its corresponding neurons
 189 "on", representing the value of that variable. Constraints are represented by
 190 inhibitory (i.e., negatively weighted) connections between the neurons. To
 191 insure that every variable is assigned a value, there is a guard neuron for

each set of neurons representing a variable; if no neuron in the set is on, the guard neuron will provide an excitatory input that is large enough to turn one on. (Because of the way the connection weights are set up, it is unlikely that the guard neuron will turn on more than one neuron.) The network is updated on each cycle by randomly picking a set of neurons that represents a variable, and flipping the state of the neuron in that set whose input is *most inconsistent* with its current output (if any). When all neurons' states are consistent with their input, a solution is achieved.

To solve the n -queens problem, for example, each of the $n \times n$ board positions is represented by a neuron whose output is either one or zero depending on whether a queen is currently placed in that position or not. (Note that this is a local representation rather than a distributed representation of the board.) If two board positions are inconsistent, then an inhibiting connection exists between the corresponding two neurons. For example, all the neurons in a column will inhibit each other, representing the constraint that two queens cannot be in the same column. For each row, there is a guard neuron connected to each of the neurons in that row which gives the neurons in the row a large excitatory input, enough so that at least one neuron in the row will turn on. The guard neurons thus enforce the constraint that one queen in each row must be on. As described above, the network is updated on each cycle by randomly picking a row and flipping the state of the neuron in that row whose input is most inconsistent with its current output. A solution is realized when the output of every neuron is consistent with its input.

5 Why does the GDS Network Perform So Well?

Our analysis of the GDS network was motivated by the following question: "Why does the network perform so much better than traditional backtracking methods on certain tasks"? In particular, we were intrigued by the results on the n -queens problem, since this problem has received considerable attention

from previous researchers. For n -queens, Adorf and Johnston found empirically that the network requires a linear number of transitions to converge. Since each transition requires linear time, the expected (empirical) time for the network to find a solution is $O(n^2)$. To check this behavior, Johnston and Adorf ran experiments with n as high as 1024, at which point memory limitations became a problem.¹

5.1 Nonsystematic Search Hypothesis

Initially, we hypothesized that the network's advantage came from the non-systematic nature of its search, as compared to the systematic organization inherent in depth-first backtracking. There are two potential problems associated with systematic depth-first search. First, the search space may be organized in such a way that poorer choices are explored first at each branch point. For instance, in the n -queens problem, depth-first search tends to find a solution more quickly when the first queen is placed in the center of the first row rather than in the corner; apparently this occurs because there are more solutions with the first queen in the center. Nevertheless, most naive algorithms tend to start in the corner simply because humans find it more natural to program that way. However, this fact by itself does not explain why nonsystematic search would work so well for n -queens. A backtracking program that randomly orders rows (and columns within rows) performs much better than the naive method, but still performs poorly relative to the GDS network.

The second potential problem with depth-first search is more significant and more subtle. As illustrated by figure 1, a depth-first search can be a disadvantage when solutions are not evenly distributed throughout the search space. In the tree at the left of the figure, the solutions are clustered together. In the tree on the right, the solutions are more evenly distributed.

¹The network, which is programmed in Lisp, requires approximately 11 minutes to solve the 1024 queens problem on a TI Explorer II. For larger problems, memory becomes a limiting factor because the network requires approximately $O(n^2)$ space. (Although the number of connections is actually $O(n^3)$, some connections are computed dynamically rather than stored).

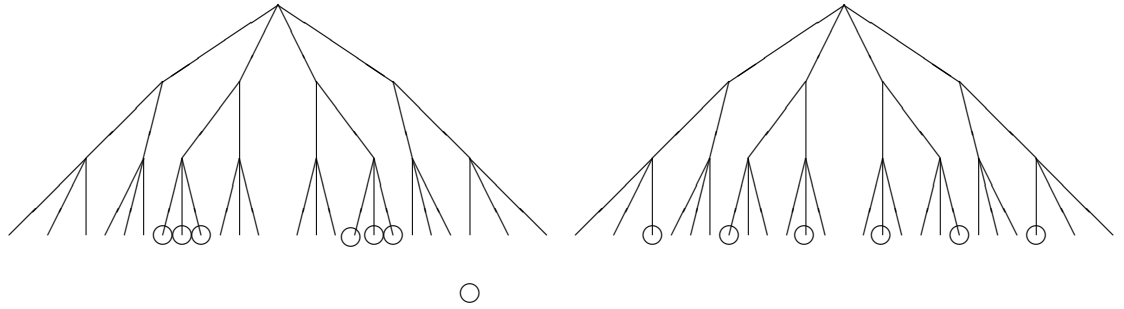


Figure 1: Solutions Clustered vs. Solutions Evenly Distributed

Thus, the average distance between solutions is greater in the left tree. In a depth-first search, the average time to find the first solution increases with the average distance between solutions. Consequently depth-first search performs relatively poorly in a tree where In comparison, a search strategy which examines the leaves of the tree in random order is unaffected by solution clustering.

We investigated whether this phenomenon explained the relatively poor performance of depth-first search on n -queens by experimenting with a randomized algorithm begins by selecting a path from the root to a leaf. To select a path, the algorithm starts at the root node and chooses one of its children with equal probability. This process continues recursively until a leaf is encountered. If the leaf is a solution the algorithm terminates, if not, it starts over again at the root and selects a path. The same path may be examined more than once, since no memory is maintained between successive trials.

The Las Vegas algorithm does, in fact, perform better than simple depth-first search on n -queens However, the performance of the Las Vegas algorithm is still not nearly as good as that of the GDS network, and so we concluded that the systematicity hypothesis alone cannot explain the network's behavior.

267 5.2 Informedness Hypothesis

268 Our second hypothesis was that the network's search process uses informa-
269 tion about the current assignment that is not available to a constructive
270 backtracking program. 's use of an iterative improvement strategy guides
271 the search in a way that is not possible with a standard backtracking algo-
272 rithm. We now believe this hypothesis is correct, in that it explains why the
273 network works so well. In particular, the key to the network's performance
274 appears to be that state transitions are made so as to reduce the number of
275 outstanding inconsistencies in the network; specifically, each state transition
276 involves flipping the neuron whose output is most inconsistent with its cur-
277 rent input. From a constraint satisfaction perspective, it is as if the network
278 reassigns a value for a variable by choosing the value that violates the fewest
279 constraints. This idea is captured by the following heuristic:

280 **Min-Conflicts heuristic:**

281 *Given:* A set of variables, a set of binary constraints, and an assign-
282 ment specifying a value for each variable. Two variables *conflict* if
283 their values violate a constraint.

284 *Procedure:* Select a variable that is in conflict, and assign it a value
285 that minimizes the number of conflicts. (Break ties randomly.)

286 We have found that the network's behavior can be approximated by a
287 symbolic system that uses the min-conflicts heuristic for hill climbing. The
288 hill-climbing system starts with an initial assignment generated in a prepro-
289 cessing phase. At each choice point, the heuristic chooses a variable that is
290 currently in conflict and reassigns its value, until a solution is found. The
291 system thus searches the space of possible assignments, favoring assignments
292 with fewer total conflicts. Of course, the hill-climbing system can become
293 "stuck" in a local maximum, in the same way that the network may become
294 "stuck" in a local minimum. In the next section we present empirical evi-
295 dence to support our claim that the min-conflicts approach can account for
296 the network's effectiveness.


```

Procedure INFORMED-BACKTRACK (VARS-LEFT VARS-DONE)
  If all variables are consistent, then solution found, STOP.
  Let VAR = a variable in VARS-LEFT that is in conflict.
  Remove VAR from VARS-LEFT.
  Push VAR onto VARS-DONE.
  Let VALUES = list of possible values for VAR in ascending order according
    to number of conflicts with variables in VARS-LEFT.
  For each VALUE in VALUES, until solution found:
    If VALUE does not conflict with any variable that is in VARS-DONE,
    then Assign VALUE to VAR.
      Call INFORMED-BACKTRACK(VARS-LEFT VARS-DONE)
    end if
  end for
end procedure

Begin program
  Let VARS-LEFT = list of all variables, each assigned an initial value.
  Let VARS-DONE = nil
  Call INFORMED-BACKTRACK(VARS-LEFT VARS-DONE)
End program

```

Figure 2: Informed Backtracking Using the Min-Conflicts Heuristic

297 There are two aspects of the min-conflicts hill-climbing method that dis-
 298 tinguish it from standard CSP algorithms. First, instead of incrementally
 299 constructing a consistent partial assignment, the min-conflicts method *re-*
 300 *pairs* a complete but inconsistent assignment by reducing inconsistencies.
 301 Thus, it uses information about the current assignment to guide its search
 302 that is not available to a standard backtracking algorithm. Second, the use
 303 of a hill-climbing strategy rather than a backtracking strategy produces a
 304 different style of search.

305 5.2.1 Repair-Based Search Strategies

306 (This is a example of a third level section.) Extracting the method from the
307 network enables us to tease apart and experiment with its different compo-
308 nents. In particular, the idea of repairing an inconsistent assignment can be
309 used with a variety of different search strategies in addition to hill climbing.
310 For example, we can backtrack through the space of possible repairs, rather
311 than using a hill-climbing strategy, as follows. Given an initial assignment
312 generated in a preprocessing phase, we can employ the min-conflicts heuristic
313 to order the choice of variables and values to consider, as described in figure
314 2. Initially, the variables are all on a list of VARS-LEFT, and as they are
315 repaired, they are pushed onto a list of VARS-DONE. The algorithm attempts
316 to find a sequence of repairs, such that no variable is repaired more than
317 once. If there is no way to repair a variable in VARS-LEFT without violat-
318 ing a previously repaired variable (a variable in VARS-DONE), the algorithm
319 backtracks.

320 Notice that this algorithm is simply a standard backtracking algorithm
321 augmented with the min-conflicts heuristic to order its choice of which vari-
322 able and value to attend to. This illustrates an important point. The back-
323 tracking repair algorithm incrementally extends a consistent partial assign-
324 ment (i.e., VARS-DONE), as does a constructive backtracking program, but
325 in addition, uses information from the initial assignment (i.e., VARS-LEFT)
326 to bias its search. Thus, it is a type of *informed backtracking*. We still char-
327 acterize it as repair-based method since its search is guided by a complete,
328 inconsistent assignment.

329 6 Experimental Results

330 [section ommitted]

331 7 A Theoretical Model

332 [section ommitted]

333 8 Discussion

334 [section ommitted]

335 9 Acknowledgement

336 The authors wish to thank Hans-Martin Adorf, Don Rosenthal, Richard
337 Franier, Peter Cheeseman and Monte Zweben for their assistance and ad-
338 vice. We also thank Ron Musick and our anonymous reviewers for their
339 comments. The Space Telescope Science Institute is operated by the Associ-
340 ation of Universities for Research in Astronomy for NASA.

341 Appendix A. Probability Distributions for N- 342 Queens

343 [section ommitted]

344 References

- 345 Cito, J., Schermann, G., Wittern, J. E., Leitner, P., Zumberi, S. and Gall,
346 H. C. (2017). An Empirical Analysis of the Docker Container Ecosystem
347 on GitHub. In *2017 IEEE/ACM 14th International Conference on Mining*
348 *Software Repositories (MSR)*, pp. 323–333, iSSN: null.
- 349 Copeland, P. (2010). Google’s Innovation Factory: Testing, Culture, and
350 Infrastructure. In *Proceedings of the 2010 Third International Conference*

351 *on Software Testing, Verification and Validation*, Washington, DC, USA:
 352 IEEE Computer Society, ICST '10, pp. 11–14.

353 Gallaba, K. and McIntosh, S. (2018). Use and Misuse of Continuous Inte-
 354 gration Features: An Empirical Study of Projects that (mis)use Travis CI.
 355 *IEEE Transactions on Software Engineering*, pp. 1–1.

356 github (2017). <https://github.blog/2017-11-07-github-welcomes-all-ci-tools/>.
 357 In github.com, ed., *github welcomes all ci tools*.

358 Kalliamvakou, E., Gousios, G., Blincoe, K., Singer, L., German, D. M. and
 359 Damian, D. (2014). The promises and perils of mining GitHub. Hyderabad,
 360 India: Association for Computing Machinery, MSR 2014, pp. 92–101.

361 Michael Hilton, K. H., Timothy Tunnell, Marinov, D. and Dig, D. (2016).
 362 Usage, costs, and benefits of continuous integration in open-source projects
 363 | Proceedings of the 31st IEEE/ACM International Conference on Auto-
 364 mated Software Engineering.

365 Rahman, A., Mahdavi-Hezaveh, R. and Williams, L. (2019). A systematic
 366 mapping study of infrastructure as code research. *Information and Soft-
 367 ware Technology*, 108, pp. 65–77.

368 Shahin, M., Ali Babar, M. and Zhu, L. (2017). Continuous Integration, De-
 369 livery and Deployment: A Systematic Review on Approaches, Tools, Chal-
 370 lenges and Practices. *IEEE Access*, 5, pp. 3909–3943.

371 Sharma, T., Fragkoulis, M. and Spinellis, D. (2016). Does Your Configuration
 372 Code Smell? In *2016 IEEE/ACM 13th Working Conference on Mining
 373 Software Repositories (MSR)*, pp. 189–200, iSSN: null.