

1 Usage and structure of continuous integra-  
2 tion as configuration?

3 Joseph Ling  
j1653@kent.ac.uk



School of Computing  
University of Kent  
United Kingdom

Word Count: 6,100

4 March 5, 2020

6 This paper describes a simple heuristic approach to solving large-scale con-  
7 straint satisfaction and scheduling problems. In this approach one starts  
8 with an inconsistent assignment for a set of variables and searches through  
9 the space of possible repairs. The search can be guided by a value-ordering  
10 heuristic, the *min-conflicts heuristic*, that attempts to minimize the num-  
11 ber of constraint violations after each step. The heuristic can be used with  
12 a variety of different search strategies. We demonstrate empirically that on  
13 the  $n$ -queens problem, a technique based on this approach performs orders of  
14 magnitude better than traditional backtracking techniques. We also describe  
15 a scheduling application where the approach has been used successfully. A  
16 theoretical analysis is presented both to explain why this method works well  
17 on certain types of problems and to predict when it is likely to be most  
18 effective.

# 1 Introduction

<https://arxiv.org/ftp/arxiv/papers/1703/1703.07019.pdf>

Continuous integration (CI) is becoming more popular over the last few years. This can be seen by how major version control hosting services Github, Bitbucket and Gitlab have all started to or have been improving their CI product. In terms of research, configuration as code Rahman, Mahdavi-Hezaveh and Williams (2019) and continuous integration Copeland (2010) with Shahin, Ali Babar and Zhu (2017) demonstrating breadth of the research.

Continuous integration is a process of automatically running compiling, running tests and checking that the product works. This can be combined with Continuous Delivery where the product is deployed or released after it has gone through CI.

This can get complicated quickly therefore configuration as code (or infrastructure as code) is used to configure it. The main kind of configuration format used for this is yaml (reference to what it is??) followed by xml and java based scripting formats.

In terms of looking at usage we are going to do a similar look at the data as did Michael Hilton, Marinov and Dig (2016). The important aspect will be looking at how usage has changed over the last 5 years along with looking more closely at which repositories are more likely to use CI/CD. For this we are going to focus on the following research questions:

- What percentage of open-source projects use CI?
- multiple CI used
- what is the breakdown of usage of different services?
- Do certain types of projects use CI more than others?

This should give us a better understanding of the sample of repositories from Github. From there we look at the structure of the configuration files to understand how certain aspects of it are used.

- 48 • configuratizon errors when loading the config (just yaml parsing errors  
49 atm)
- 50 • how are comments used in the configuration?
- 51 • how scripts with the configuration files? (need to elloborate more on  
52 this one)

53 A key aspect is that these questions do not look too deeply into the  
54 individual implementation of each CI system. This is because there are  
55 already some good papers looking Gallaba and McIntosh (2018) at this but  
56 in order to be able to compare the different configuration types it is important  
57 to compare similar attributes (there is also a time factor in here as well).

## 58 2 Previous Works

### 59 2.1 Continous integration

60 Continous integration is the frequent submission of work normally tied into  
61 a feedback loop. For example using version control daily committing changes.  
62 That then a server builds and tests the changes informing you of status of  
63 those cahnges. The generally agree upon detailed definition is Fowler (2010).

### 64 2.2 Usage of continous integration

65 The actual usage of continous integration as configuration was looked at  
66 by Michael Hilton, Marinov and Dig (2016). In this they use three source  
67 of information github repositories, travis builds and a survery. In order to  
68 be do a more systematic study of CI usage than Vasilescu et al. (2015). In  
69 analysing that data they found that "The trends that we discovered point  
70 to an expected growth of CI. In the future, CI will have an even greater  
71 influence than it has today.". As we are looking at the same question we will  
72 use four of the same research questions out of the fourteen. In order to see  
73 what difference four years has made to the growth of usage of CI.

## 74 2.3 Config as code

75 Configuration as code or infrastructure as code has been an increasing area  
76 of research over the last few years. There seems to be slightly more research  
77 in infrastructure as code Rahman, Mahdavi-Hezaveh and Williams (2019).  
78 There has been a focus on Puppet and Chef, for example in Sharma, Frangkoulis  
79 and Spinellis (2016) looks at code quality by the measure of "code smell" of  
80 Puppet code. This tackles the problem by defining by best practices and  
81 analyzing the code against that. In the case of Cito et al. (2017) it uses  
82 the docker linter in order to be able to analyse the files. For the continuous  
83 integration systems we pick we will look into the tooling around that to aid  
84 the analysis.

## 85 3 Methodology

86 In order to get repositories with CI/CD configuration from Github we have a  
87 number of approaches. The first is to use the search for particular files but  
88 this is limited to only 1000 results. The alternative is to search for repositories  
89 and we bypass the 1000 result limit to an extent by getting results for every  
90 'star' count (stars are used to like or upvote a repository). Although this will  
91 be giving us a lot of results it will still only be a sample of the population but  
92 will give us a wider range of results. As there is rate limiting multiple github  
93 api keys can be used to speed up the scraping of data (github could also  
94 be used to speed up the process I think).

95 After we have got a repository we need to get the CI/CD files from it.  
96 This is fairly easy as the CI/CD systems normally require a strict naming  
97 convention and location within the repository. However as most of them are  
98 yaml based you can have ".yaml" and ".yml" and users can use all sorts of  
99 mixtures of upper and lower case. We try to account for this but won't get  
100 every scenario. This combined with the fact that we are only looking for  
101 top configuration files based on github (2017) along with github actions and  
102 azure pipelines. Is why we also check repositories for their README.md file

103 to check if it has a build tag.

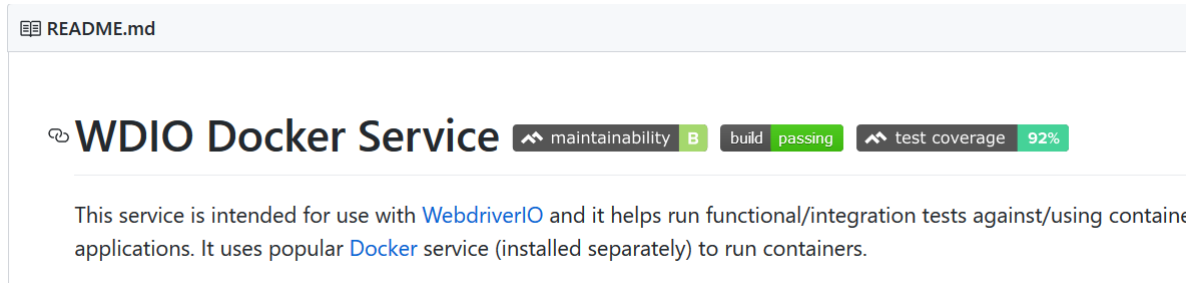


Figure 1: Example of CI tag for Github ReadMe

where did this image come from??  
reference it man

104  
105 In doing so it should give a wider net when sampling and help to under-  
106 stand when a CI system is either not using configuration as code or using a  
107 different CI system.

108 There are dangers in scraping data off github in terms of assumptions to  
109 do with the population as found in Kalliamvakou et al. (2014). Our dataset  
110 does not contain any forked repositories. But due to time constraints number  
111 of commits and frequency of recent commits has not been looked at. This  
112 would be an interesting area of further research in order to improve the  
113 quality of the sample but also to look at how that affects the frequency of  
114 CI usage.

115 Additionally the assumption that all repositories are of programming  
116 projects with code in them is wrong. A number of repositories can be used  
117 for storage, experimental, academic and other things. However they to all  
118 some extent can use CI/CD for their work as a number of books were found  
119 when looking through the dataset could use CI/CD.

120 Tooling for the configuration files, I looked into Travis, Github Actions  
121 and Jenkins to work out whether or not it could aid in the research or not.  
122 As a key part of understanding the first relies on knowing whether or not it is  
123 valid. In terms for travis there is currently two parsers to validate the config-  
124 uration. One which is deprecated since 2017 travis (2017) the other which is  
125 currently in development travis (2020). Both didn't provided the necessary  
126 results with the most recent one not being able to handle default fields. For

127 Github Actions as it's still a new tooling for it hasn't been developed out-  
 128 side of the Github editor web page (<https://github.community/t5/GitHub-Actions/YAML-validator-for-Github-Actions-possible-expansion-of/td-p/29557>).  
 129 For Jenkins which is older solution allows validation through http/ssh request  
 130 to the Jenkins server (Gitlab follows this style as well) Jenkins (2020) Gitlab  
 131 (2020). This could work well although would require setting up a server for  
 132 each configuration type and might not validate if variables from the config  
 133 aren't defined on the server. As well as it would be best to be able to validate  
 134 them all or none of them in terms of being able to compare results easily.  
 135

## 136 4 Usage of CI

### 137 4.1 What percentage of open-source projects use CI?

138 Based a search for configuration as configuration files for the following CI  
 139 systems: Travis, Gitlab, Azure, App Veyor, Drone, Jenkins, Github, Circleci,  
 140 Semaphore, Teamcity and buildkite. Wrecker got bought by Oracle and from  
 141 doing a search on Github for what I think based on the docs (docs: Wrecker  
 142 and Oracle (2018) and search: GitHub (2020)) for their config file naming  
 143 convention. I was only able to find 20 results so did not include in the scraping  
 144 script to speed up the process of searching for the other configuration file  
 145 formats.

CI/CD	count	repos with config	no. multiple	multiple percent
config file(s)	12128	38.51%	1675	13.81%
found in ReadMe	873	2.77%		
none found	18493	58.72%		

Table 1: Percentage of CI used for projects

146 Our sample of repositories is 31,494 in comparison to Michael Hilton,  
 147 Marinov and Dig (2016) which had a sample of 34,544. The percentage of  
 148 CI projects they had was 40.27%. As if you combined the "config file(s)"

149 and "found in ReadMe". However in order to work out if a project might be  
150 using CI but the config file wasn't picked a search string is used. Therefore  
151 it is not as accurate as finding a config file as their could be false postives.

152 However that doesn't give us too much insight into the dataset. Here is a  
153 graph showing the subscribers plotted against the number of stars. The key  
154 here to understand is not potentially any correlation but to see the spread  
155 of data that the table is showing.




Figure 2: Scatter graph of Github stars against subscribers

156 Figure 2 helps give a understanding to the give a depth of the data for  
157 where the graph is just blue. This is because on Github you get more repos-  
158 itories with smaller star counts than large ones.

159 Figure 3 provides insight into the density of the data for between 0 to  
160 25000.



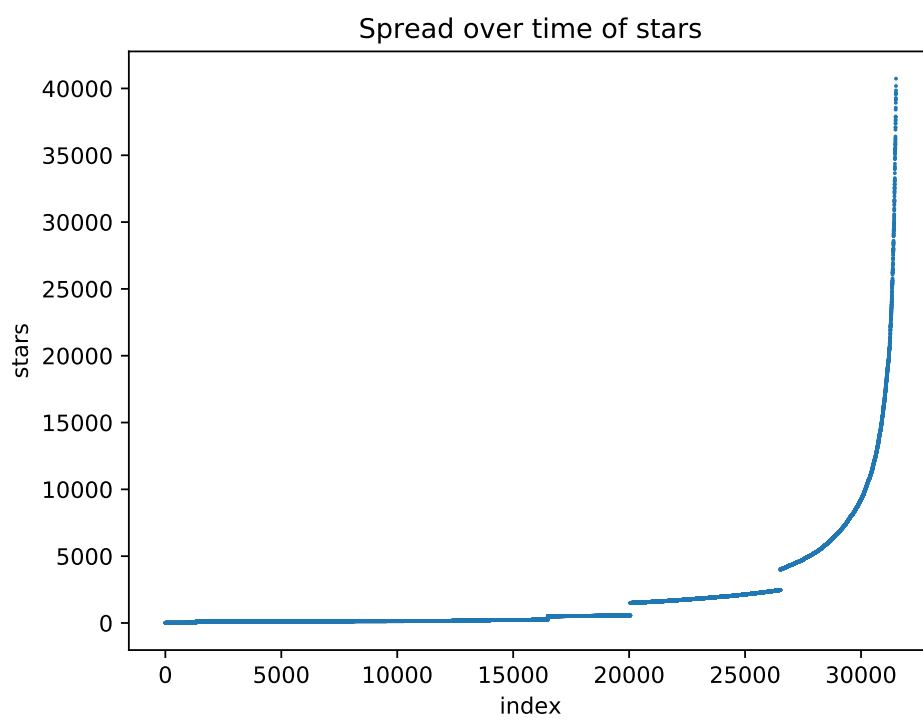


Figure 3: Stars graph

## 161 4.2 What CI systems are projects using?

162 Like all other research travis is the most popular CI system in use. However  
163 over the last 4 years since the github (2017) Circleci has lost out on it's rough  
164 quarter that it owned. In particular the rise of github actions seems to have  
165 taken second place even though it is still very young in comparison (DATES).  
166 However this might not be down to the Circleci loosing out on their existing  
167 share. But potentially as the rise in CI usage goes up on github. Projects  
are more likely to pick in the built in solutions to github.

Table 2: dogs

	config	percentage
travis	10273	74%
github	2190	16%
circleci	1066	8%
jenkinsPipeline	151	1%
drone	83	1%
buildkite	32	0%
teamcity	4	0%
semaphore	2	0%
azure	1	0%

168

169 naming convention????

## 170 4.3 Do certain types of projects use CI more than oth- 171 ers?

172 Looking back on Michael Hilton, Marinov and Dig (2016) results (GET THE  
173 GRAPH IN).

174 They observed that in the sampled groups that had more stars had a  
175 higher percentage of projects using CI. In particular that the group with the  
176 highest star count had a 70% of the projects using CI.

177 Figure 4 shows all the CI projects sorted then grouped together per 540




Figure 4: Stars graph

178 of them. Then the average of whether or not that a project for that group  
 179 use CI is displayed on the graph.

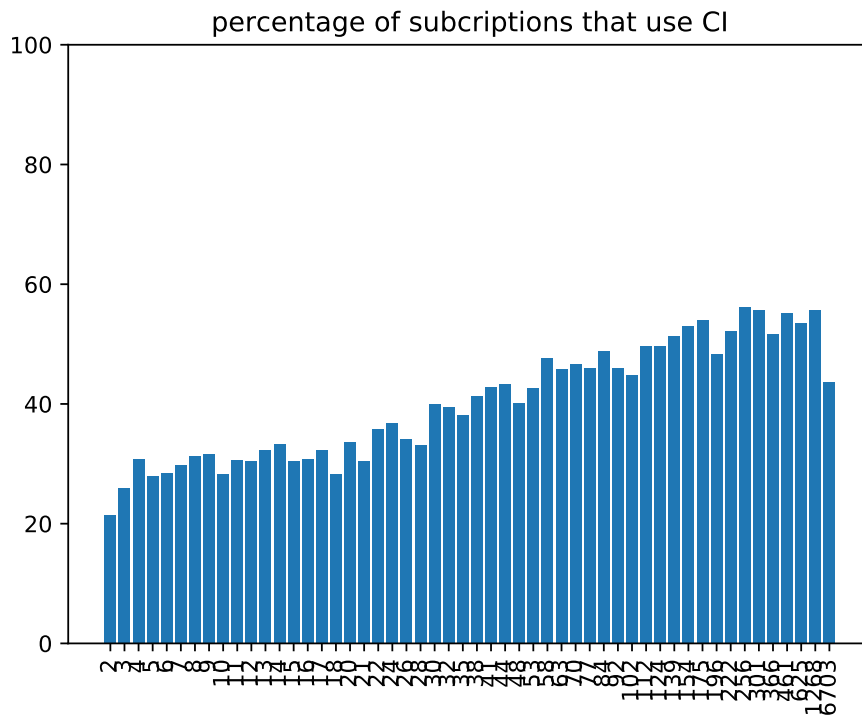


Figure 5: Stars graph

180 Figure 5 uses the same method as Figure 4 except is does it based the  
 181 number of subscribers. Subscribers are used on github to keep update on  
 182 the changes on the project. This ranges from core team memembers working  
 183 on the project to people that want to be notified about a new release. In  
 184 looking at this metric the hypothesis was that it would have a sharper rise  
 185 in percentage of projects using CI per subscriber. However that was not the  
 186 case overall the gradient is not as strong along with the fact that the final  
 187 group dramatically dips (note: check sample size of that group).

188 After looking at these results three areas of research that could be inter-  
 189 esting would be filtering the sample to only projects that had recent commits.  
 190 On the assumption that CI is becoming more popular and stale projects won't

191 require CI.  
 192 Looking into whether or not that assumption is the case. Then looking  
 193 commit count to see if projects with more commits are more likely to use CI.

## 194 5 Config file results

### 195 5.1 configuration errors when loading the config (just 196 yaml parsing errors atm)

Table 3: cats

yaml_encoding_error	composer error	parse error	scanner error
config			
circleci	0	0	1
drone	30	0	0
github	0	0	3
travis	6	10	21

### 197 5.2 How are comments used in configuration?

### 198 5.3 How are script tags used?

199 - how scripts with the configuration files? (need to elaborate more on this  
 200 one)

201 By almost any measure, the Hubble Space Telescope scheduling problem  
 202 Between ten thousand and thirty thousand astronomical observations per  
 203 year must be scheduled, subject to a great variety of constraints including  
 204 power restrictions, observation priorities, time-dependent orbital character-  
 205 istics, movement of astronomical bodies, stray light sources, etc. Because the  
 206 telescope is an extremely valuable resource with a limited lifetime, efficient  
 207 scheduling is a critical concern. An initial scheduling system, developed us-  
 208 ing traditional programming methods, highlighted the difficulty of the prob-  
 209 lem; it was estimated that it would take over three weeks for the system to

## Analysis of yaml continous integration file formats




Figure 6: Sankey daigram

210 schedule one week of observations. As described in section 7, this problem  
 211 was remedied by the development of a successful constraint-based system to  
 212 augment the initial system. At the heart of the constraint-based system is  
 213 a neural network developed by Adorf and Johnston, the Guarded Discrete  
 214 Stochastic (GDS) network,

215 From a computational point of view the network is interesting because  
 216 Adorf and Johnston found that it performs well on a variety of tasks, in  
 217 addition to the space telescope scheduling problem. For example, the network  
 218 performs significantly better on the  $n$ -queens problem than methods that  
 219 were previously developed. The  $n$ -queens problem requires placing  $n$  queens  
 220 on an  $n \times n$  chessboard so that no two queens share a row, column or diagonal.  
 221 The network has been used to solve problems of up to 1024 queens, whereas  
 222 most heuristic backtracking methods encounter difficulties with problems  
 223 one-tenth

224 In a standard Hopfield network, all connections between neurons are sym-  
 225 metric. In the GDS network, the main network is coupled asymmetrically to  
 226 an auxiliary network of *guard neurons* which restricts the configurations that  
 227 the network can assume. This modification enables the network to rapidly  
 228 find a solution for many problems, even when the network is simulated on  
 229 a serial machine. Unfortunately, convergence to a stable configuration is no  
 230 longer guaranteed. Thus the network can fall into a local minimum involving  
 231 a group of unstable states among which it will oscillate. In practice, however,  
 232 if the network fails to converge after some number of neuron state transitions,  
 233 it can simply be stopped and started over.

234 To illustrate the network architecture and updating scheme, let us con-  
 235 sider how the network is used to solve binary constraint satisfaction problems.  
 236 A problem consists of  $n$  variables,  $X_1 \dots X_n$ , with domains  $D_1 \dots D_n$ , and a  
 237 set of binary constraints. Each constraint  $C_\alpha(X_j, X_k)$  is a subset of  $D_j \times D_k$   
 238 specifying incompatible values for a pair of variables. The goal is to find  
 239 an assignment for each of the variables which satisfies the constraints. (In  
 240 this paper we only consider the task of finding a single solution, rather than  
 241 that of finding all solutions.) To solve a CSP using the network, each vari-  
 242 able is represented by a separate set of neurons, one neuron for each of the  
 243 variable's possible values. Each neuron is either "on" or "off", and in a solu-  
 244 tion state, every variable will have exactly one of its corresponding neurons  
 245 "on", representing the value of that variable. Constraints are represented by  
 246 inhibitory (i.e., negatively weighted) connections between the neurons. To  
 247 insure that every variable is assigned a value, there is a guard neuron for  
 248 each set of neurons representing a variable; if no neuron in the set is on, the  
 249 guard neuron will provide an excitatory input that is large enough to turn  
 250 one on. (Because of the way the connection weights are set up, it is unlikely  
 251 that the guard neuron will turn on more than one neuron.) The network is  
 252 updated on each cycle by randomly picking a set of neurons that represents  
 253 a variable, and flipping the state of the neuron in that set whose input is  
 254 *most inconsistent* with its current output (if any). When all neurons' states

255 are consistent with their input, a solution is achieved.

256 To solve the  $n$ -queens problem, for example, each of the  $n \times n$  board posi-  
257 tions is represented by a neuron whose output is either one or zero depending  
258 on whether a queen is currently placed in that position or not. (Note that  
259 this is a local representation rather than a distributed representation of the  
260 board.) If two board positions are inconsistent, then an inhibiting connection  
261 exists between the corresponding two neurons. For example, all the neurons  
262 in a column will inhibit each other, representing the constraint that two  
263 queens cannot be in the same column. For each row, there is a guard neuron  
264 connected to each of the neurons in that row which gives the neurons in the  
265 row a large excitatory input, enough so that at least one neuron in the row  
266 will turn on. The guard neurons thus enforce the constraint that one queen  
267 in each row must be on. As described above, the network is updated on each  
268 cycle by randomly picking a row and flipping the state of the neuron in that  
269 row whose input is most inconsistent with its current output. A solution is  
270 realized when the output of every neuron is consistent with its input.

## 271 6 Why does the GDS Network Perform So 272 Well?

273 Our analysis of the GDS network was motivated by the following question:  
274 “Why does the network perform so much better than traditional backtracking  
275 methods on certain tasks”? In particular, we were intrigued by the results on  
276 the  $n$ -queens problem, since this problem has received considerable attention  
277 from previous researchers. For  $n$ -queens, Adorf and Johnston found empir-  
278 ically that the network requires a linear number of transitions to converge.  
279 Since each transition requires linear time, the expected (empirical) time for  
280 the network to find a solution is  $O(n^2)$ . To check this behavior, Johnston  
281 and Adorf ran experiments with  $n$  as high as 1024, at which point memory



282 limitations became a problem.<sup>1</sup>

## 283 6.1 Nonsystematic Search Hypothesis

284 Initially, we hypothesized that the network's advantage came from the non-  
285 systematic nature of its search, as compared to the systematic organization  
286 inherent in depth-first backtracking. There are two potential problems as-  
287 sociated with systematic depth-first search. First, the search space may be  
288 organized in such a way that poorer choices are explored first at each branch  
289 point. For instance, in the  $n$ -queens problem, depth-first search tends to find  
290 a solution more quickly when the first queen is placed in the center of the  
291 first row rather than in the corner; apparently this occurs because there are  
292 more solutions with the queen in the center. Nevertheless, most naive algorithms tend to start  
293 in the corner simply because humans find it more natural to program that  
294 way. However, this fact by itself does not explain why nonsystematic search  
295 would work so well for  $n$ -queens. A backtracking program that randomly  
296 orders rows (and columns within rows) performs much better than the naive  
297 method, but still performs poorly relative to the GDS network.

298 The second potential problem with depth-first search is more significant  
299 and more subtle. As illustrated by figure 7, a depth-first search can be  
300 a disadvantage when solutions are not evenly distributed throughout the  
301 search space. In the tree at the left of the figure, the solutions are clustered  
302 together. In the tree on the right, the solutions are more evenly distributed.  
303 Thus, the average distance between solutions is greater in the left tree. In a  
304 depth-first search, the average time to find the first solution increases with the  
305 average distance between solutions. Consequently depth-first search performs  
306 relatively poorly in a tree where solutions are clustered. In comparison, a search strategy which  
307 examines the leaves of the tree in random order is unaffected by solution

---

<sup>1</sup>The network, which is programmed in Lisp, requires approximately 11 minutes to solve the 1024 queens problem on a TI Explorer II. For larger problems, memory becomes a limiting factor because the network requires approximately  $O(n^2)$  space. (Although the number of connections is actually  $O(n^3)$ , some connections are computed dynamically rather than stored).




Figure 7: Solutions Clustered vs. Solutions Evenly Distributed

clustering.

We investigated whether this phenomenon explained the relatively poor performance of depth-first search on  $n$ -queens by experimenting with a randomized algorithm begins by selecting a path from the root to a leaf. To select a path, the algorithm starts at the root node and chooses one of its children with equal probability. This process continues recursively until a leaf is encountered. If the leaf is a solution the algorithm terminates, if not, it starts over again at the root and selects a path. The same path may be examined more than once, since no memory is maintained between successive trials.

The Las Vegas algorithm does, in fact, perform better than simple depth-first search on  $n$ -queens. However, the performance of the Las Vegas algorithm is still not nearly as good as that of the GDS network, and so we concluded that the systematicity hypothesis alone cannot explain the network's behavior.

## 6.2 Informedness Hypothesis

Our second hypothesis was that the network's search process uses information about the current assignment that is not available to a constructive backtracking program. 's use of an iterative improvement strategy guides the search in a way that is not possible with a standard backtracking algo-

328 rithm. We now believe this hypothesis is correct, in that it explains why the  
329 network works so well. In particular, the key to the network's performance  
330 appears to be that state transitions are made so as to reduce the number of  
331 outstanding inconsistencies in the network; specifically, each state transition  
332 involves flipping the neuron whose output is most inconsistent with its cur-  
333 rent input. From a constraint satisfaction perspective, it is as if the network  
334 reassigns a value for a variable by choosing the value that violates the fewest  
335 constraints. This idea is captured by the following heuristic:

336       **Min-Conflicts heuristic:**

337       *Given:* A set of variables, a set of binary constraints, and an assign-  
338       ment specifying a value for each variable. Two variables *conflict* if  
339       their values violate a constraint.

340       *Procedure:* Select a variable that is in conflict, and assign it a value  
341       that minimizes the number of conflicts. (Break ties randomly.)

342       We have found that the network's behavior can be approximated by a  
343       symbolic system that uses the min-conflicts heuristic for hill climbing. The  
344       hill-climbing system starts with an initial assignment generated in a prepro-  
345       cessing phase. At each choice point, the heuristic chooses a variable that is  
346       currently in conflict and reassigns its value, until a solution is found. The  
347       system thus searches the space of possible assignments, favoring assignments  
348       with fewer total conflicts. Of course, the hill-climbing system can become  
349       “stuck” in a local maximum, in the same way that the network may become  
350       “stuck” in a local minimum. In the next section we present empirical evi-  
351       dence to support our claim that the min-conflicts approach can account for  
352       the network's effectiveness.

353       There are two aspects of the min-conflicts hill-climbing method that dis-  
354       tinguish it from standard CSP algorithms. First, instead of incrementally  
355       constructing a consistent partial assignment, the min-conflicts method *re-*  
356       *pairs* a complete but inconsistent assignment by reducing inconsistencies.  
357       Thus, it uses information about the current assignment to guide its search  
358       that is not available to a standard backtracking algorithm. Second, the use

```

Procedure INFORMED-BACKTRACK (VARS-LEFT VARS-DONE)
  If all variables are consistent, then solution found, STOP.
  Let VAR = a variable in VARS-LEFT that is in conflict.
  Remove VAR from VARS-LEFT.
  Push VAR onto VARS-DONE.
  Let VALUES = list of possible values for VAR in ascending order according
                  to number of conflicts with variables in VARS-LEFT.
  For each VALUE in VALUES, until solution found:
    If VALUE does not conflict with any variable that is in VARS-DONE,
      then Assign VALUE to VAR.
      Call INFORMED-BACKTRACK(VARS-LEFT VARS-DONE)
    end if
  end for
end procedure

Begin program
  Let VARS-LEFT = list of all variables, each assigned an initial value.
  Let VARS-DONE = nil
  Call INFORMED-BACKTRACK(VARS-LEFT VARS-DONE)
End program

```

Figure 8: Informed Backtracking Using the Min-Conflicts Heuristic

359 of a hill-climbing strategy rather than a backtracking strategy produces a  
 360 different style of search.

### 361 6.2.1 Repair-Based Search Strategies

362 (This is an example of a third level section.) Extracting the method from the  
 363 network enables us to tease apart and experiment with its different compo-  
 364 nents. In particular, the idea of repairing an inconsistent assignment can be  
 365 used with a variety of different search strategies in addition to hill climbing.

366 For example, we can backtrack through the space of possible repairs, rather  
367 than using a hill-climbing strategy, as follows. Given an initial assignment  
368 generated in a preprocessing phase, we can employ the min-conflicts heuristic  
369 to order the choice of variables and values to consider, as described in figure  
370 8. Initially, the variables are all on a list of VARS-LEFT, and as they are  
371 repaired, they are pushed onto a list of VARS-DONE. The algorithm attempts  
372 to find a sequence of repairs, such that no variable is repaired more than  
373 once. If there is no way to repair a variable in VARS-LEFT without violat-  
374 ing a previously repaired variable (a variable in VARS-DONE), the algorithm  
375 backtracks.

376 Notice that this algorithm is simply a standard backtracking algorithm  
377 augmented with the min-conflicts heuristic to order its choice of which vari-  
378 able and value to attend to. This illustrates an important point. The back-  
379 tracking repair algorithm incrementally extends a consistent partial assign-  
380 ment (i.e., VARS-DONE), as does a constructive backtracking program, but  
381 in addition, uses information from the initial assignment (i.e., VARS-LEFT)  
382 to bias its search. Thus, it is a type of *informed backtracking*. We still char-  
383 acterize it as repair-based method since its search is guided by a complete,  
384 inconsistent assignment.

## 385 7 Experimental Results

386 [section ommitted]

## 387 8 A Theoretical Model

388 [section ommitted]

## 389 9 Discussion

390 [section ommitted]

## 391 10 Acknowledgement

392 The authors wish to thank Hans-Martin Adorf, Don Rosenthal, Richard  
393 Franier, Peter Cheeseman and Monte Zweben for their assistance and ad-  
394 vice. We also thank Ron Musick and our anonymous reviewers for their  
395 comments. The Space Telescope Science Institute is operated by the Associ-  
396 ation of Universities for Research in Astronomy for NASA.

## 397 Appendix A. Probability Distributions for N- 398 Queens

399 [section omitted]

## 400 References

- 401 Cito, J., Schermann, G., Wittern, J. E., Leitner, P., Zumberi, S. and Gall,  
402 H. C. (2017). An Empirical Analysis of the Docker Container Ecosystem  
403 on GitHub. In *2017 IEEE/ACM 14th International Conference on Mining*  
404 *Software Repositories (MSR)*, pp. 323–333, iSSN: null.
- 405 Copeland, P. (2010). Google’s Innovation Factory: Testing, Culture, and  
406 Infrastructure. In *Proceedings of the 2010 Third International Conference*  
407 *on Software Testing, Verification and Validation*, Washington, DC, USA:  
408 IEEE Computer Society, ICST ’10, pp. 11–14.
- 409 Fowler, M. (2010). Continuous integration. In  
410 <https://www.martinfowler.com/articles/continuousIntegration.html>.
- 411 Gallaba, K. and McIntosh, S. (2018). Use and Misuse of Continuous Inte-  
412 gration Features: An Empirical Study of Projects that (mis)use Travis CI.  
413 *IEEE Transactions on Software Engineering*, pp. 1–1.

414 github (2017). <https://github.blog/2017-11-07-github-welcomes-all-ci-tools/>.  
415 In github.com, ed., *github welcomes all ci tools*.

416 GitHub (2020). github filename search for wrecker.yml files. In *github file-*  
417 *name search for wrecker.yml files*.

418 Gitlab (2020). <https://docs.gitlab.com/ee/api/lint.html>. In *Gitlab docs*.

419 Jenkins (2020). <https://jenkins.io/doc/book/pipeline/development/>. In  
420 *Jenkins documentation*.

421 Kalliamvakou, E., Gousios, G., Blincoe, K., Singer, L., German, D. M. and  
422 Damian, D. (2014). The promises and perils of mining GitHub. Hyderabad,  
423 India: Association for Computing Machinery, MSR 2014, pp. 92–101.

424 Michael Hilton, K. H., Timothy Tunnell, Marinov, D. and Dig, D. (2016).  
425 Usage, costs, and benefits of continuous integration in open-source projects  
426 | Proceedings of the 31st IEEE/ACM International Conference on Auto-  
427 mated Software Engineering.

428 Rahman, A., Mahdavi-Hezaveh, R. and Williams, L. (2019). A systematic  
429 mapping study of infrastructure as code research. *Information and Soft-*  
430 *ware Technology*, 108, pp. 65–77.

431 Shahin, M., Ali Babar, M. and Zhu, L. (2017). Continuous Integration, De-  
432 livery and Deployment: A Systematic Review on Approaches, Tools, Chal-  
433 lenges and Practices. *IEEE Access*, 5, pp. 3909–3943.

434 Sharma, T., Fragkoulis, M. and Spinellis, D. (2016). Does Your Configuration  
435 Code Smell? In *2016 IEEE/ACM 13th Working Conference on Mining*  
436 *Software Repositories (MSR)*, pp. 189–200, iSSN: null.

437 travis (2017). travis yaml (old repository). In [https://github.com/travis-](https://github.com/travis-ci/travis-yaml/)  
438 [ci/travis-yaml/](https://github.com/travis-ci/travis-yaml/).

439 travis (2020). travis yaml new implementation. In [https://github.com/travis-](https://github.com/travis-ci/travis-yaml/)  
440 [ci/travis-yaml/](https://github.com/travis-ci/travis-yaml/).

- 441 Vasilescu, B., Yu, Y., Wang, H., Devanbu, P. and Filkov, V. (2015). Quality  
442 and productivity outcomes relating to continuous integration in GitHub.  
443 Bergamo, Italy: Association for Computing Machinery, ESEC/FSE 2015,  
444 pp. 805–816.
- 445 Wrecker and Oracle (2018). Wrecker ci development blog. In *Wrecker CI*  
446 *development blog*.